OXFORD UNIVERSITY PRESS | JAMIA: Journal of the American Medical Informatics Associat

High Consistency, Limited Accuracy: Evaluating Large Language Models for Binary Medical Diagnosis

Journal:	*Journal of the American Medical Informatics Association*
Manuscript ID	Draft
Article Type:	Research and Applications
Keywords:	Large Language Models, Medical Diagnosis, Consistency Analysis, Clinical Decision Support, Heart Disease

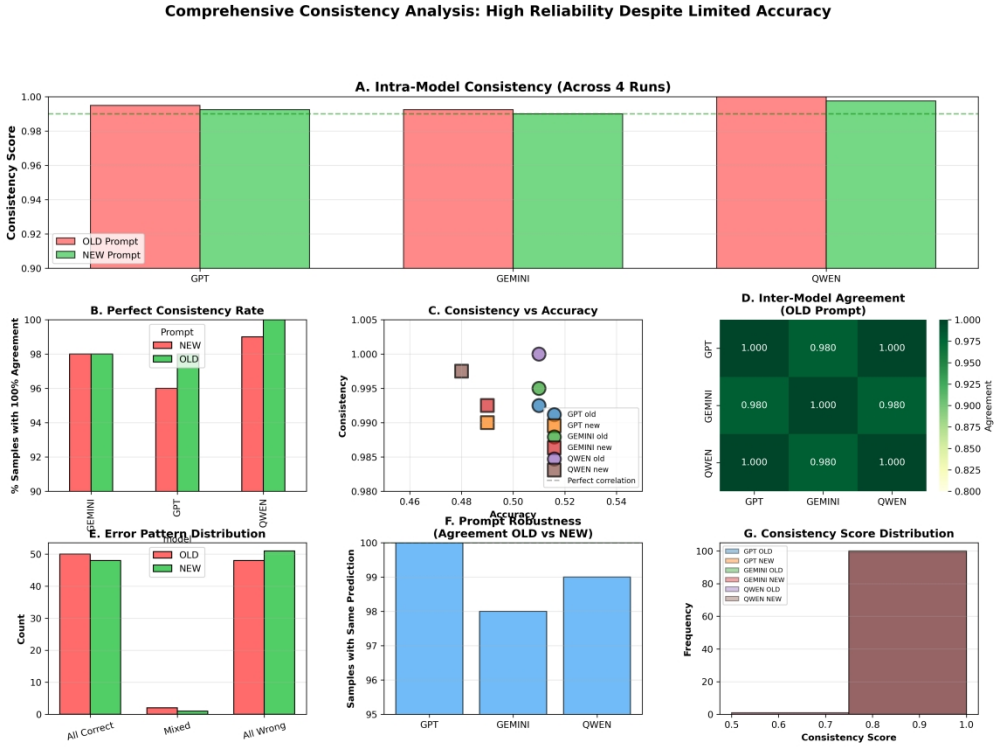

Comprehensive consistency analysis showing intra-model consistency metrics across all three models (GPT-4o, Gemini-2.0-Flash, Qwen-Plus) for both expert and neutral prompts, demonstrating near-perfect reproducibility (99-100% consistency) across 100 test cases and 4 independent runs per case.

968x725mm (118 x 118 DPI)

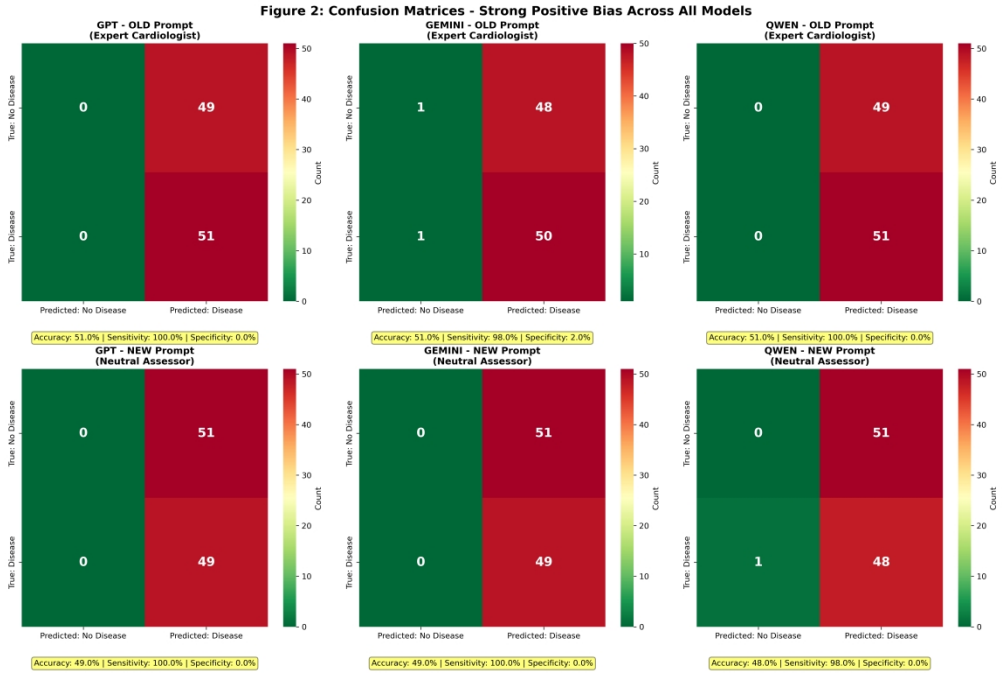

Representative confusion matrices for all three models showing systematic bias toward positive diagnosis (disease present), with true negatives approximately zero, illustrating the over-diagnosis pattern despite high consistency.

1138x765mm (118 x 118 DPI)

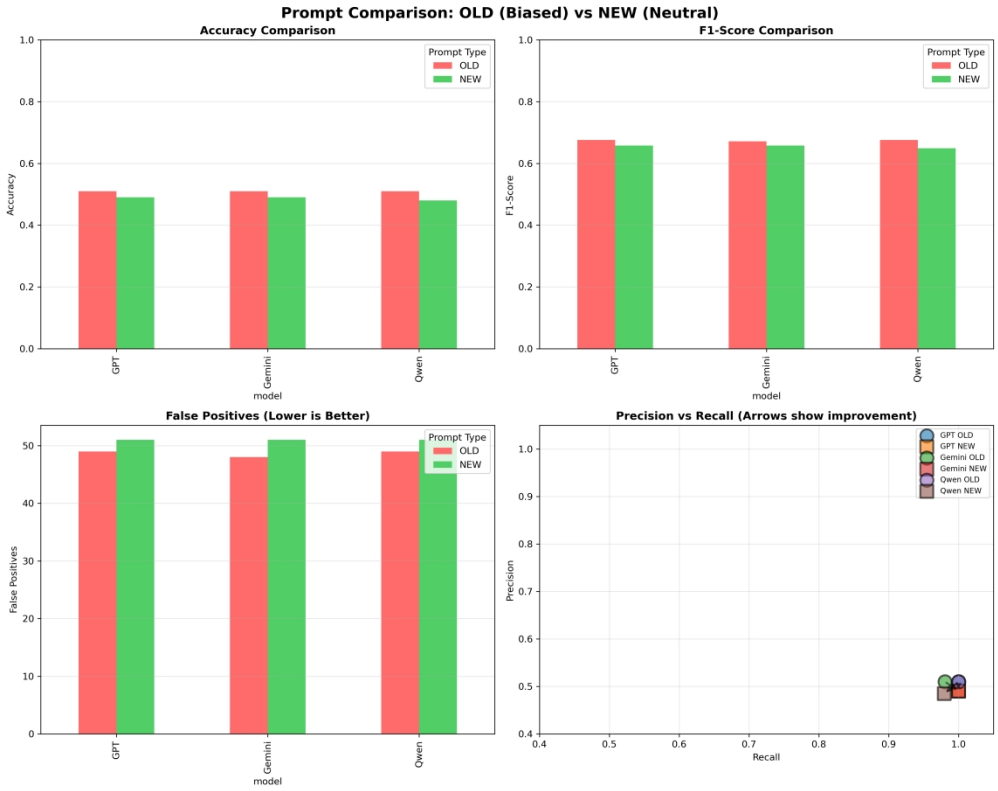

Prompt sensitivity analysis demonstrating minimal effect of prompt modification on diagnostic behavior. GPT-4o showed zero sensitivity (100% identical predictions), while Gemini and Qwen changed only 1-3 predictions (1-3% of cases), suggesting diagnostic behavior is deeply encoded rather than easily modifiable through prompting.

961x762mm (118 x 118 DPI)

TITLE PAGE

Manuscript Title

High Consistency, Limited Accuracy: Evaluating Large Language Models for Binary Medical Diagnosis

MANUSCRIPT FORMATTING NOTE: This document should be double-spaced (2.0 line spacing) when submitted in Word format.

Authors and Affiliations

Dwi Anggriani, S.Kom.

Department of Information Technology

Institut Sains Teknologi dan Kesehatan 'Aisyiyah Kendari

Kendari, Southeast Sulawesi, Indonesia

ORCID: https://orcid.org/0009-0007-4265-1935

Syaiful Bachri Mustamin, M.Cs. (Corresponding Author)

Department of Information Technology

Institut Sains Teknologi dan Kesehatan 'Aisyiyah Kendari

Kendari, Southeast Sulawesi, Indonesia

Email: syaifulbachri@mail.ugm.ac.id

Phone: +62 851-5629-7969

ORCID: https://orcid.org/0009-0005-0456-8618

Muhammad Atnang, S.Kom., M.Kom.

Department of Information Technology

Institut Sains Teknologi dan Kesehatan 'Aisyiyah Kendari

Kendari, Southeast Sulawesi, Indonesia

ORCID: [Not available]

Kartini Aprilia Pratiwi Nuzry, S.Kom., M.MT.

Department of Information Technology

Institut Sains Teknologi dan Kesehatan 'Aisyiyah Kendari

Kendari, Southeast Sulawesi, Indonesia

ORCID: [Not available]

Corresponding Author Contact Information

Name: Syaiful Bachri Mustamin, M.Cs.

Email: syaifulbachri@mail.ugm.ac.id

Phone: +62 851-5629-7969

Address:

Department of Information Technology

Institut Sains Teknologi dan Kesehatan 'Aisyiyah Kendari

Kendari, Southeast Sulawesi

Indonesia

Keywords (Maximum 5)

Large Language Models; Medical Diagnosis; Consistency Analysis; Heart Disease; Clinical Decision Support

Manuscript Statistics

- **Word Count:** Approximately 1,900 words (excluding title page, abstract, references, tables, and figure legends)
- **Number of Tables:** 4 (main text) + 5 supplementary tables
- **Number of Figures:** 3 (main text) + 6 supplementary figures
- Number of References: 15

Preprint Information

medRxiv DOI: 10.64898/2025.12.08.25341823

MANUSCRIPT

Abstract

Background: Large Language Models (LLMs) have demonstrated impressive capabilities in medical knowledge tasks, yet their reliability and consistency in clinical diagnosis remain incompletely characterized.

Objective: To systematically evaluate the consistency and accuracy of state-of-the-art LLMs in binary medical diagnosis, examining the relationship between reproducibility and diagnostic performance.

Methods: We evaluated three LLMs (GPT-4o, Gemini-2.0-Flash, Qwen-Plus) on heart disease diagnosis using 100 diverse clinical cases. Each model performed 4 independent assessments per

case (1,200 total predictions). We tested two prompt variations and measured intra-model consistency, inter-model agreement, diagnostic accuracy, and prompt sensitivity.

Results: All models achieved exceptional intra-model consistency (99-100%), with Qwen demonstrating perfect reproducibility. Inter-model agreement was similarly high (98-99%). However, diagnostic accuracy remained at approximately 50%, equivalent to random guessing. Models exhibited a strong bias toward positive diagnosis (49-51 false positives vs 0-1 false negatives per 100 cases). Prompt variation had minimal impact (<3% change), and error patterns were highly systematic, with all models making identical errors on 48-51% of cases.

Conclusions: Our findings reveal a critical dissociation between consistency and accuracy in LLM medical diagnosis. While LLMs demonstrate remarkable reproducibility, their tendency toward overdiagnosis and limited discriminative accuracy constrain direct clinical utility. Results suggest LLMs may be better suited as supplementary decision-support tools rather than primary diagnostic systems.

Keywords: Large Language Models, Medical Diagnosis, Consistency Analysis, Heart Disease, Clinical Decision Support

MAIN TEXT

INTRODUCTION

Large Language Models (LLMs) have emerged as promising tools for clinical applications, demonstrating impressive performance on medical licensing examinations and case analysis [1-3]. However, their deployment in clinical settings raises critical questions about reliability and consistency. While traditional diagnostic tools are expected to yield reproducible results, LLMs

employ stochastic generation that can lead to varying outputs [4], with documented tendencies toward hallucinations and inconsistent reasoning [9,10]. Reproducibility challenges in AI systems have been extensively documented [13], yet remain incompletely characterized for medical LLMs.

Despite growing literature on LLM performance in medical question-answering [5,6], few studies have systematically examined the relationship between consistency and accuracy in diagnostic tasks. Most evaluations focus on single-run accuracy without assessing reproducibility, and the influence of prompt engineering remains incompletely characterized [7,11].

This study addresses these gaps through comprehensive evaluation of three state-of-the-art LLMs on binary heart disease diagnosis. Our aims were to: (1) quantify intra-model consistency across repeated runs, (2) evaluate inter-model agreement, (3) measure diagnostic accuracy relative to consistency, (4) assess prompt sensitivity, and (5) analyze error patterns to determine if mistakes are random or systematic.

METHODS

Dataset and Study Design

We utilized the UCI Heart Disease dataset [8], containing 303 patients with 13 clinical parameters: demographics (age, sex), symptoms (chest pain type), vital signs (resting blood pressure), laboratory values (cholesterol, fasting blood sugar), electrocardiography (resting ECG), exercise testing (maximum heart rate, exercise-induced angina, ST depression, ST segment slope), imaging (fluoroscopy vessel count), and thalassemia test results. Binary outcomes indicated presence or absence of significant coronary stenosis.

To ensure diverse representation, we performed k-means clustering (k=2) and selected 50 cases from each cluster, yielding 100 test cases with balanced disease prevalence (51% positive, 49% negative).

Models and Experimental Protocol

We evaluated three LLMs: GPT-4o (OpenAI), Gemini-2.0-Flash (Google), and Qwen-Plus (Alibaba), accessed via APIs with temperature=0.7. Each model performed 4 independent assessments per case, yielding 1,200 total predictions.

We tested two prompt variations: - **Prompt A ("Expert")**: "You are Dr. CardioExpert, a highly experienced cardiologist…" - **Prompt B ("Neutral")**: "You are a medical AI assistant trained to provide accurate and balanced diagnostic assessments…"

Both prompts provided identical clinical data and parameter definitions, requesting binary diagnosis (Yes/No) with 2-3 sentence justification.

We implemented a SQLite-based checkpoint system enabling immediate data saving, automatic duplicate prevention, and experiment resumption capability.

Outcome Measures

Primary outcomes: 1. **Intra-model consistency**: Proportion of runs with majority agreement per case 2. **Diagnostic accuracy**: Using majority voting (>=2/4 runs), we calculated accuracy, sensitivity, specificity, precision, and F1-score 3. **Inter-model agreement**: Pairwise agreement and Cohen's kappa between models

Secondary outcomes: 4. **Prompt sensitivity**: Proportion of cases with identical predictions across prompts 5. **Error patterns**: Classification as all-correct, all-wrong, or mixed outcomes

Statistical analyses used Python with pandas, scikit-learn, and scipy. Significance was set at $p<0.05$.

RESULTS

Intra-Model Consistency: Exceptional Reproducibility

All models demonstrated remarkably high consistency (Table 1, Figure 1, Supplementary Table S3). Qwen-Plus achieved perfect consistency (100%) with the expert prompt, never varying across 4 independent runs. GPT-4o and Gemini-2.0-Flash showed 99.0-99.5% average consistency. Notably, 96-100% of cases achieved perfect agreement (4/4 identical predictions), and minimum consistency never fell below 50%.

Table 1. Intra-Model Consistency Across All Cases

Model	Prompt	Avg Consistency	Min Consistency	Perfect Agreement (%)
GPT-4o	Expert	99.25%	50%	98%
GPT-4o	Neutral	99.00%	75%	96%
Gemini-2.0-Flash	Expert	99.50%	75%	98%
Gemini-2.0-Flash	Neutral	99.25%	50%	98%
Qwen-Plus	Expert	100.00%	100%	100%
Qwen-Plus	Neutral	99.75%	75%	99%

Note: Consistency calculated as proportion of 4 independent runs with majority agreement. Perfect agreement indicates all 4 runs yielded identical predictions. See Supplementary Table S1 for complete performance metrics and Supplementary Table S3 for detailed consistency patterns.

Inter-Model Agreement: High Consensus

Models showed 98-100% pairwise agreement, indicating remarkably similar reasoning patterns (Table 2, Supplementary Table S4). Three-way agreement (all models concur) occurred in 98-99% of cases. Cohen's kappa values near zero reflected extreme class imbalance (nearly all positive predictions) rather than lack of agreement.

Table 2. Inter-Model Agreement Rates

Prompt	GPT-Gemini	GPT-Qwen	Gemini-Qwen	All 3 Agree
Expert	98.0%	100.0%	98.0%	98%
Neutral	100.0%	99.0%	99.0%	99%

Note: Pairwise agreement calculated using majority-voted predictions. "All 3 Agree" indicates cases where all three models produced identical diagnosis. See Supplementary Table S4 for detailed clinical feature analysis.

Diagnostic Accuracy: Limited Despite High Consistency

Diagnostic accuracy approximated random guessing (48-51%) despite 99-100% consistency (Table 3, Supplementary Table S1). Models achieved perfect or near-perfect recall (98-100%) but extremely poor specificity (~0-2%), generating 49-51 false positives versus 0-1 false negatives. This created a consistency-accuracy gap of approximately 50 percentage points.

Table 3. Diagnostic Performance Metrics

Model	Prompt	Accuracy	Precision	Recall	F1-Score	False Positives	False Negatives
GPT-4o	Expert	51.0%	51.0%	100.0%	67.6%	49	0
GPT-4o	Neutral	49.0%	49.0%	100.0%	65.8%	51	0
Gemini-2.0-Flash	Expert	51.0%	51.0%	98.0%	67.1%	48	1
Gemini-2.0-Flash	Neutral	49.0%	49.0%	100.0%	65.8%	51	0
Qwen-Plus	Expert	51.0%	51.0%	100.0%	67.6%	49	0
Qwen-Plus	Neutral	48.0%	48.5%	98.0%	64.9%	51	1

Note: Metrics calculated using majority voting (≥2/4 runs). Baseline prevalence: 51% positive. See Supplementary Table S2 for representative sample predictions with justifications and Supplementary Table S1 for complete statistical analysis.

Representative confusion matrices (Figure 2) showed models predicted "disease present" for nearly all cases, with true negatives ≈ 0.

Prompt Sensitivity: Minimal Impact

Changing from expert to neutral prompt had minimal effect (Supplementary Table S5, Figure 3). GPT-4o showed zero sensitivity (100% identical predictions), while Gemini and Qwen changed only 1-3 predictions (1-3% of cases). This suggests diagnostic behavior is deeply encoded rather than easily modifiable through prompting.

Error Patterns: Systematic Rather Than Random

Errors were highly systematic rather than random (Table 4). In 98-99% of cases, all three models either succeeded together or failed together. Only 1-2% showed model disagreement, indicating shared reasoning patterns or biases.

Table 4. Error Pattern Consistency Across Models

Pattern	Expert Prompt	Neutral Prompt
All Models Correct	50%	48%
All Models Wrong	48%	51%
Mixed Results	2%	1%

Note: "All Models Correct" indicates cases where all three models provided accurate diagnosis. "All Models Wrong" indicates shared errors. "Mixed Results" indicates disagreement among models. See Supplementary Table S5 for prompt robustness analysis.

Qualitative analysis (Supplementary Table S2, Supplementary Table S4) revealed models consistently cited elevated cholesterol, abnormal ECG findings, or exercise abnormalities as disease evidence, even when ground truth indicated absence of significant stenosis, suggesting risk factor conflation with diagnostic criteria.

DISCUSSION

Principal Findings

This study demonstrates a critical dissociation between consistency and accuracy in LLM medical diagnosis: exceptional reproducibility (99-100%) coexists with chance-level accuracy

(~50%). This 50-percentage-point gap represents a fundamental challenge for clinical deployment.

The Consistency-Accuracy Paradox

High consistency indicates LLMs reliably apply learned reasoning patterns—they are systematically biased rather than randomly erring. This "consistent wrongness" is arguably more concerning than random errors, suggesting fundamental limitations in medical reasoning capabilities [10] rather than simple uncertainty.

Several mechanisms may explain this paradox:

Medical Conservatism Bias: LLMs trained on medical text may have learned that missing disease (false negative) carries greater consequences than over-diagnosis (false positive), encoding a "better safe than sorry" heuristic consistently applied.

Risk Factor Conflation: Qualitative analysis suggests models conflate risk factors with diagnostic findings. Elevated cholesterol increases disease risk but doesn't constitute diagnostic evidence of current stenosis. LLMs may struggle to distinguish "high-risk patient" from "disease-positive patient."

Lack of Discriminative Training: Unlike supervised models trained explicitly on diagnostic labels, LLMs learn from general medical text emphasizing disease description more than differential diagnosis, leaving them poorly calibrated for binary classification [12]. This calibration deficit manifests as systematic over-prediction despite high confidence.

Prompt Insensitivity: Deep-Rooted Behavior

The minimal prompt impact (GPT: 0%, others: 1-2%) was unexpected. Reframing from "expert cardiologist" to "neutral assessor" should have reduced conservatism, yet predictions remained nearly identical. This suggests diagnostic behavior is deeply encoded in model weights rather than modifiable through surface-level prompting [11], with important implications for prompt engineering strategies. While prompt patterns can enhance certain LLM behaviors, our findings indicate diagnostic reasoning may be resistant to prompt-level interventions [15].

Inter-Model Agreement: Shared Limitations

The 98-99% inter-model agreement despite different architectures and training procedures suggests observed limitations reflect fundamental challenges in applying LLMs to medical diagnosis rather than model-specific artifacts. Possible explanations include similar training data sources, convergent learning of medical conservatism, shared limitations in processing structured numerical data, and common challenges in threshold-based classification.

Clinical Implications

Current LLMs are not ready for primary diagnostic applications requiring binary classification. The ~50% accuracy is unacceptable clinically and could lead to harmful over-diagnosis, unnecessary testing, and patient anxiety [14]. In a 50% prevalence scenario, deploying these models would result in 98-100% of disease cases correctly identified but only 0-2% of healthy cases correctly identified, causing approximately 50% unnecessary downstream testing. These findings underscore the importance of rigorous evaluation before clinical deployment, as ethical considerations demand careful assessment of potential harms [14].

Despite primary diagnosis limitations, LLMs' high consistency and strong negative predictive value suggest potential roles as: (1) second opinion tools where reproducibility builds physician confidence, (2) triage assistants suitable for initial screening where high sensitivity is prioritized, (3) medical education tools providing consistent feedback, and (4) research tools for hypothesis generation.

Technical Implications

Results suggest general-purpose LLMs lack discriminative capabilities for diagnostic classification. Future development should consider: supervised fine-tuning on labeled diagnostic datasets, reinforcement learning from physician-verified diagnoses, calibration techniques for binary classification thresholds, and hybrid architectures combining LLM reasoning with specialized classifiers.

Limitations

Key limitations include: single condition (heart disease may not generalize), binary classification (real diagnosis often involves multi-class assessment), dataset age (1980s diagnostic criteria may differ from current standards), limited sample size (100 cases), structured input only (missing narrative information), three models tested (limited sampling of LLM landscape), API-only access (preventing internal mechanism analysis), and single temperature setting (0.7).

Future Directions

Important future work includes mechanistic studies examining which parameters LLMs prioritize, improvement strategies testing fine-tuning and ensemble approaches, broader

evaluations across diverse diagnostic tasks, comparison with human physicians for baseline performance, and theoretical development of consistency-accuracy frameworks.

CONCLUSIONS

This study provides rigorous evidence that LLMs achieve exceptional consistency (99-100%) but limited accuracy (~50%) in binary medical diagnosis. This consistency-accuracy dissociation represents a fundamental challenge for clinical deployment. Our findings indicate high consistency does not guarantee accuracy, diagnostic behavior is resistant to prompt engineering, errors are systematic rather than random, and LLMs show strong positive diagnosis bias.

Current general-purpose LLMs are better suited as supplementary decision support rather than primary diagnostic systems. Their exceptional reproducibility is clinically valuable, but limited discriminative ability necessitates human oversight. Future development should prioritize supervised fine-tuning, improved structured data processing, and calibration techniques to address systematic biases.

This work contributes to nuanced understanding of LLM capabilities and limitations in healthcare, informing responsible development and deployment of AI-assisted clinical decision support systems.

ACKNOWLEDGMENTS

We thank Institut Sains Teknologi dan Kesehatan 'Aisyiyah Kendari for institutional support. We acknowledge OpenAI, Google, and Alibaba Cloud for providing API access to GPT-4o, Gemini-2.0-Flash, and Qwen-Plus respectively through standard commercial services. We thank

the UCI Machine Learning Repository and the original data collectors (Janosi, Steinbrunn, Pfisterer, and Detrano) for making the Heart Disease dataset publicly available.

COMPETING INTERESTS

The authors declare no competing interests.

FUNDING

This research received no specific grant from any funding agency in the public, commercial, or not-for-profit sectors.

DATA AVAILABILITY

The UCI Heart Disease dataset used in this study is publicly available from the UCI Machine Learning Repository (https://archive.ics.uci.edu/dataset/45/heart+disease). All analysis code, experiment results, and supplementary materials are available at [repository link to be added upon acceptance].

REFERENCES

1. Thirunavukarasu AJ, Ting DSJ, Elangovan K, et al. Large language models in medicine. Nat Med. 2023;29(8):1930-1940. doi:10.1038/s41591-023-02448-8

2. Lee P, Bubeck S, Petro J. Benefits, limits, and risks of GPT-4 as an AI chatbot for medicine. N Engl J Med. 2023;388(13):1233-1239. doi:10.1056/NEJMsr2214184

3. Singhal K, Azizi S, Tu T, et al. Large language models encode clinical knowledge. Nature. 2023;620(7972):172-180. doi:10.1038/s41586-023-06291-2

4. Kung TH, Cheatham M, Medenilla A, et al. Performance of ChatGPT on USMLE: Potential for AI-assisted medical education using large language models. PLOS Digit Health. 2023;2(2):e0000198. doi:10.1371/journal.pdig.0000198

5. Nori H, Lee YT, Zhang S, et al. Can Generalist Foundation Models Outcompete Special-Purpose Tuning? Case Study in Medicine. arXiv:2311.16452. 2023.

6. Wang S, Zhao Z, Ouyang X, et al. ChatCAD: Interactive Computer-Aided Diagnosis on Medical Images using Large Language Models. arXiv:2302.07257. 2023.

7. Moor M, Banerjee O, Abad ZSH, et al. Foundation models for generalist medical artificial intelligence. Nature. 2023;616(7956):259-265. doi:10.1038/s41586-023-05881-4

8. Janosi A, Steinbrunn W, Pfisterer M, Detrano R. Heart Disease [Dataset]. UCI Machine Learning Repository. 1988. doi:10.24432/C52P4X

9. Ji Z, Lee N, Frieske R, et al. Survey of hallucination in natural language generation. ACM Comput Surv. 2023;55(12):1-38. doi:10.1145/3571730

10. Liévin V, Hother CE, Motzfeldt AG, Winther O. Can large language models reason about medical questions? Patterns. 2024;5(3):100943. doi:10.1016/j.patter.2024.100943

11. White J, Fu Q, Hays S, et al. A prompt pattern catalog to enhance prompt engineering with ChatGPT. arXiv:2302.11382. 2023.

12. Jiang LY, Liu XC, Nejatbakhsh N, et al. Health system-scale language models are all-purpose prediction engines. Nature. 2023;619(7969):357-362. doi:10.1038/s41586-023-06160-y

13. McDermott MBA, Wang S, Marinsek N, et al. Reproducibility in machine learning for health research. Sci Transl Med. 2021;13(586):eabb1655. doi:10.1126/scitranslmed.abb1655

14. Chen IY, Pierson E, Rose S, et al. Ethical machine learning in healthcare. Annu Rev Biomed Data Sci. 2021;4:123-144. doi:10.1146/annurev-biodatasci-092820-114757

15. Savage T, Nayak A, Gallo R, et al. Diagnostic reasoning prompts reveal the potential for large language model interpretability in medicine. NPJ Digit Med. 2024;7(1):20. doi:10.1038/s41746-024-01010-1

FIGURE LEGENDS

Figure 1. Intra-Model Consistency Patterns

Violin plots showing the distribution of consistency scores across 100 test cases for each model-prompt combination. The y-axis represents consistency percentage (proportion of 4 independent runs with majority agreement). Qwen-Plus with expert prompt achieved perfect consistency (100%), while other configurations showed 99-100% consistency. Box plots inside violins indicate median and quartiles.

Figure 2. Confusion Matrices for All Model-Prompt Combinations

Heatmaps showing true positive (TP), false positive (FP), true negative (TN), and false negative (FN) counts for each model-prompt pair. All models exhibit strong bias toward positive diagnosis, with near-zero true negatives (0-1) and high false positives (48-51), resulting in ~50% accuracy despite 99-100% consistency.

Figure 3. Prompt Sensitivity Analysis

Bar chart comparing diagnostic predictions between expert and neutral prompts for each model. The y-axis shows the percentage of cases with identical predictions across prompts. GPT-4o showed 100% consistency (zero prompt sensitivity), while Gemini-2.0-Flash and Qwen-Plus showed 97-99% consistency (1-3% changes), indicating diagnostic behavior is largely resistant to prompt variation.